MiniBase Content Moderation Agent

Week 2 FDE Program Deliverable - Scenario 4

The Problem: Moderation at Capacity

MiniBase's hybrid moderation team is drowning in routine spam, leaving no capacity for nuanced judgment cases.

Metric	Current State	Impact
Daily posts flagged	1,500/day	12.5% of all posts
Team capacity	47 hrs/day	8 volunteers + 2 staff
Routine spam	1,080/day (72%)	9 hrs/day consumed
Grey-zone cases	360/day (24%)	30 hrs/day needed
Burnout risk	HIGH	Volunteers quitting

The Stakes: "False positives are survivable; one viral false negative is existential."

Why NOT "Just Automate Everything"?

The Trap: Assuming all moderation work is the same ignores asymmetric risk.

- **Routine spam (72% of queue):** High volume, low risk, clear patterns → Automate
- **Grey-zone cases (24% of queue):** Lower volume, existential risk → Requires human judgment

Grey-zone examples requiring cultural context:

- Harsh critique vs. harassment (tone, subforum norms, invitation)
- Community member vs. commercial spam (established vs. new user)
- Sponsor sensitivities (the "2024 incident" precedent)

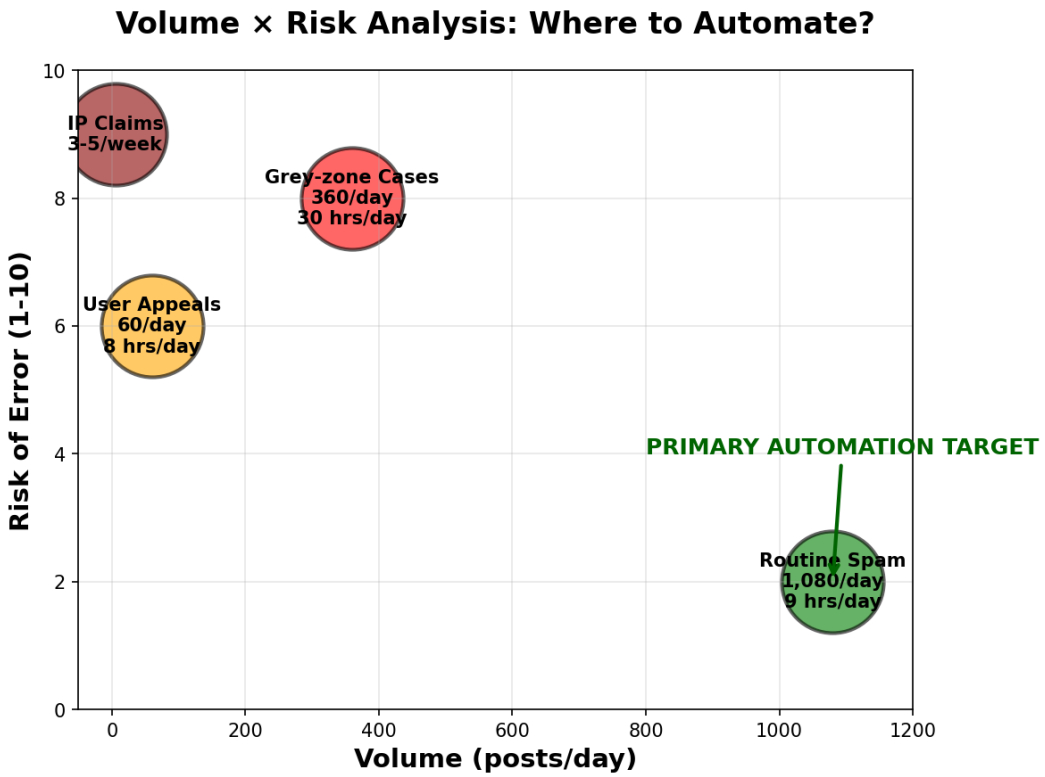

The Insight: Automate 9 hrs/day of low-risk spam to create capacity for 30 hrs/day of high-risk judgment.

The Solution: Four Delegation Archetypes

Not one-size-fits-all automation — right delegation for each task type.

Archetype	Volume/Day	Use Case	Value
■ Fully Agentic	800 posts	Clear spam (link farms, gibberish)	7 hrs saved
■ Agent-Led	280 posts	Edge cases → propose, human approves	3 hrs saved
■ Human-Led	360 posts	Grey-zone → agent gathers context	2 hrs saved
■ Human-Only	31 posts	IP claims, bans → no agent role	0 hrs saved

Safety Bar (Always Escalate):

- Sponsor accounts (2024 incident precedent) → Tom reviews personally
- High community engagement (>12 reactions) → risky to auto-remove
- Low confidence (<0.7) → requires human judgment
- Subforum-specific norms (e.g., painters "no critique without invitation")

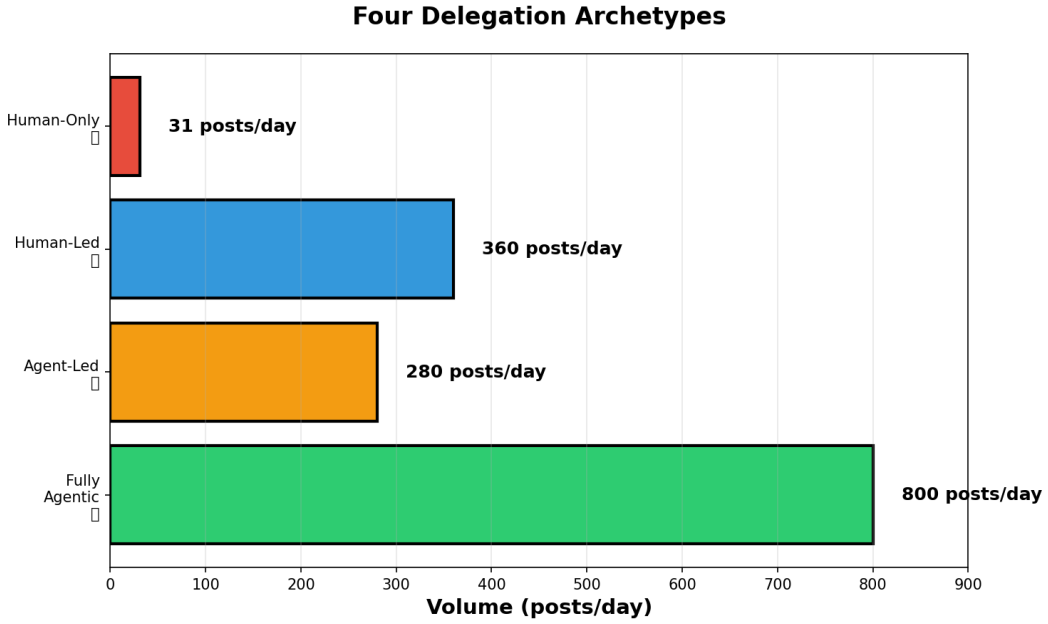

The Value & Safety Case

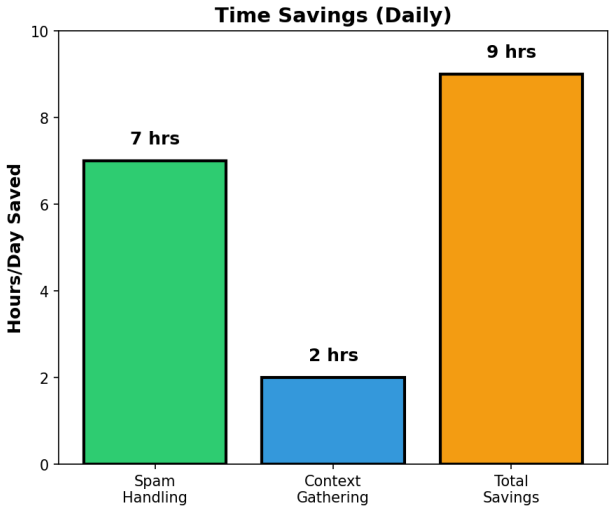

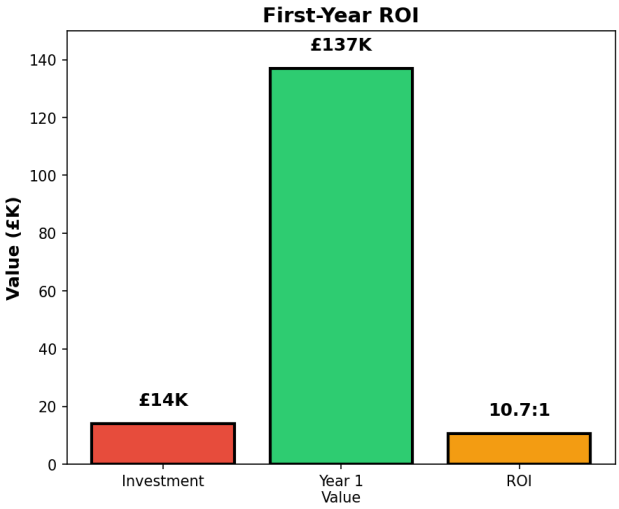

Value Metrics	Target	Safety Metrics (Non-Negotiable)	Target
Time saved on spam	77% reduction (9→2 hrs/day)	Moderation accuracy	≥95% precision
Capacity freed	7 hrs/day for grey-zone cases	False negative rate	<0.1% (zero viral)
Growth headroom	33% increase (1,500→2,000/day)	Sponsor safety	100% escalation
First-year ROI	10.7:1 (£137K / £14K)	Audit trail	100% logged

Key Message: This isn't about replacing moderators — it's about freeing them from brain-dead spam work so they can focus on nuanced cases that build community trust.

Validation: Proven Buildable

What We Built (Week 2 FDE Deliverables):

- **Agent Purpose Document:** 18 micro-tasks mapped to delegation archetypes, confidence thresholds, escalation triggers
- **System Integration Plan:** 8 systems analyzed, 212 hrs effort with specific technical constraints addressed (no hand-waving)
- **Discovery Questions:** 18 design-changing questions for Tom (6 critical, 9 high priority)
- **Live Demo:** StreamLit app validating delegation logic on 4 test cases (link farm spam, harsh critique, user appeal, IP claim)

Demo Validation Results:

- Link farm spam → Auto-remove (Fully Agentic, confidence 0.98)
- Harsh critique → Escalate (Human-Led, confidence 0.65, subforum norm conflict)
- User appeal → Context gathered (Human-Led, 5 minutes saved)
- IP claim → Human-only (No agent role, legal sensitivity)

Phase	Timeline	Key Activities	Success Criteria
Discovery	Week 3	Answer 18 critical questions with Tom	All design-changing questions resolved
Phase 1: Build	Weeks 4-19 (16 weeks)	Build spam automation + escalation logic Integrate 8 systems (Discourse, Gallery, etc.)	System tests pass ≥95% accuracy on test fixtures
Shadow Mode	Weeks 20-23 (4 weeks)	Run parallel with humans Calibrate confidence thresholds	Zero false negatives Moderator trust established
Go-Live	Week 24	Phase 1 production deployment	7 hrs/day savings achieved <10% override rate

Closed Build Loop Outcome: Specification is complete and buildable. Ready to proceed to discovery and implementation.